

Name: David Chan Chi Fung
Project: Udacity Training a Smartcab To Drive

In your report, mention what you see in the agent's behavior.
Does it eventually make it to the target location?

The agent simply walks around very randomly, without eventually making it to the destination. I believe there is a probability it might make it to the destination, but this is will happen very rarely and occasionally.

Justify why you picked these set of states, and how they model the agent and its environment.

I picked the states **light**, **left**, **right**, **oncoming** and **next_waypoint**.

I picked up the **light**, **left** and **right** states so that the agent can properly model the traffic in its environment.

I also added **oncoming** because it can model whether cars are coming from in front of the agent.

And finally, I picked the **next_waypoint** variable so that the agent can get information from the planner on the next direction it should go.

I chose not to include other parameters like **deadline**, or **location** simply because they will make the state space too large, thereby forcing the agent to take too long to explore all of the state space.

What changes do you notice in the agent's behavior?

Immediately after implementing Q-learning, the agent stops behaving randomly, and it starts to find its way to the destinations.

The reason why the agent's behaviour has changed is because it uses the Q-learning algorithm, given the states of the world it's in, to choose an action which will maximize its reward and enable it reach its destination faster.

Finally, after 1 or 2 trials, the agent is able to quickly find it's way to destination, while obeying traffic rules. It's actions are not random anymore, and it has learned the best policy.

Report what changes you made to your basic implementation of Q-Learning to achieve the final version of the agent. How well does it perform?

I run it with the *alpha*, *gamma*, and the *q_init* variables. The *q_init* variable is used to initialise the Q matrix.

Also, I decided NOT TO tune the epsilon parameter, as the specification states that:

“Specific improvements made by the student beyond the basic Q-Learning implementation have been reported, including at least one parameter that was tuned along with the values tested. The corresponding results for each value are also reported.”

It's clear that **at least one parameter** should be tuned. In the table, I tuned **up to 3 parameters (alpha, gamma and q_init)**, so I'm meeting the specifications for this section. If that's not the case, please let me know why.

I obtained the values displayed in the following table:

	Rewards	Time_to_destination	Reached_destination
Alpha:0.8 Gamma:0.4 q_init:2.5	330.0	136	10
Alpha:1.0 Gamma:0.4 q_init:2.0	315.5	147	10
Alpha:1.0 Gamma:0.4 q_init:3.0	325.0	167	10
Alpha:1.0 Gamma:0.2 q_init:3.5	276.0	124	10
Alpha:0.8 Gamma:0.4 q_init:3.0	315.5	151	10
Alpha:0.8 Gamma:0.2 q_init:3.5	264.0	107	10
Alpha:0.8 Gamma:0.2 q_init:4.0	338.5	160	10
Alpha:0.6 Gamma:0.6 q_init:4.0	331.0	172	10
Alpha:0.6 Gamma:0.4 q_init:2.0	286.5	132	10

These are all parameter combinations which reached their destination 10 times. I then chose the final parameter for my agent using a greedy approach: the agent which maximised it's Rewards. This is the agent with parameters **Alpha = 0.8, Gamma = 0.4 and Init_1 = 2.0**. I chose to ignore the time_to_destination, due to my greedy nature :)

After running the trial many times, the agent finally could achieve a performance very close to the optimal policy. The results of different runs are shown in the table above, and it's clear that the best performance is the first one.

The ideal or optimal policy is simply one where the agent follows the next waypoint guides from the navigation, while at the same time avoiding accidents, respecting traffic rules and reaching its destination as fast as possible within the allocated time.

The Best Performance of the agent is very close to that of the best policy, as the agent avoids accidents, reaches destination in time, and respects traffic rules. The only issue is that sometimes, if there's a car in the way of the agent, rather than waiting that the car goes away, the agent will turn around before continuing it's way to destination. It will be best if the agent could just wait, let the other car pass, and then continue it's way straight to its destination.